

THE THALES PLAYBOOK

Knowledge, markets, and the timing of informed trades — a paper-trading research ledger

I • Why Thales

The project is named for Thales of Miletus, the philosopher whom Aristotle credits with the first recorded act of informed speculation. Studying the winter sky, Thales concluded that the coming olive harvest would be exceptional. He did not buy olives, and he did not buy groves. He quietly rented every olive press in Miletus and Chios — off-season, when nobody wanted them and deposits were cheap — and when the bumper harvest arrived and everyone needed pressing capacity at once, he let the presses out on his own terms.

Three things make the story a founding charter rather than a parable. First, the edge was **knowledge**: a verifiable inference about the physical world, not a rumour about prices. Second, the expression was **structural**: he bought the bottleneck (the presses), not the commodity — the tollbooth, not the harvest. Third, and decisively, the profit came from **timing**: he acted while the information was cheap and the market indifferent, long before the harvest made the insight obvious to everyone.

The Thales project asks whether that triangle — knowledge, market structure, and timing — can be run as a disciplined, falsifiable research process today. Everything in this playbook is executed on paper. The ledger's product is not profit; it is calibration: written hypotheses, pre-registered kill criteria, and graded predictions.



Thales of Miletus, c. 585 BC — the first recorded option on future capacity

Fig. 1 — The founding triangle: an informed inference, expressed through market structure, at the right moment.

II • Hypotheses first, tickers second

Every play in the ledger begins the same way: a **verifiable hypothesis is postulated in writing before any market data is pulled**. Only then is the stock market scoured for tickers likely to move positively in line with the prediction, barring major upsets. The order matters. Writing the thesis first prevents the oldest failure mode in research — finding a chart that already went up and inventing a story to fit it.

Discipline is enforced by **pre-registered kill criteria**, executed by the data rather than by mood: **K1 viability** (adequate market cap and liquidity; no going-concern candidates), **K2 seat authenticity** (the asset must genuinely be what the thesis needs it to be — conglomerates trading on an unrelated narrative die here), **K3 window survival** (positive trailing-year return and drawdown better than roughly -60% ; an asset that round-tripped -75% failed its holders whatever its moat), **K4 self-funding** (no forced dilutive raises inside the window), and **K5 non-redundancy** (pairwise correlation below ~ 0.9 against a seat already filled — one bet must not wear two hats).

Surviving candidates face an **anchor-point backtest**: search-verified price anchors across the trailing year, interpolated monthly, with drawdowns explicitly declared as understated versus daily reality. Every number needs a source; nothing is recalled from memory. The backtest is treated as evidence, not oracle — a kill may stand against a winning backtest if the forward mechanism justifies it, but that argument must be made out loud. Each run ends in a dated, falsifiable log entry with scheduled checkpoints, and an anti-goalpost rule: kill criteria may be amended only *between* runs, never mid-run.




Fig. 2 — The hypothesis funnel: thesis → scour → gauntlet → backtest → dated log entry.

III · Pairs, and what pairs taught us

The first structural step beyond single-name bets was the **pair**: two assets exposed to the same hypothesis from different angles, so that the book expresses the thesis rather than a single company's execution risk. Pairs were considered wherever a thesis had both a growth expression and a defensive or structural expression — the harvest and the press, in Thales's terms.

Pairs proved instructive mainly by their failures. The recurring flaw was **redundancy**: two names that look like diversification but are one bet wearing two hats. The canonical lesson in the ledger is silver against platinum — a candidate that returned +67% over the window yet was killed at K5 on a 0.95 monthly correlation to a seat already occupied. Holding both would have doubled exposure while pretending to spread it. The opposite flaw also appeared: pairs so unrelated that no single hypothesis connected them, which makes attribution impossible — when the book moves, you cannot say whether the thesis or the noise did it.

A useful pair, we concluded, needs a stated **relationship** between its legs: complementary seats in one thesis, or a spread whose convergence the hypothesis predicts. That conclusion pushed the project in two directions at once — toward pairs with an explicit arbitrage-like relationship reviewed across sectors, and toward a third seat that insures the regime the first two share. Both are covered overleaf.

Honesty note for reviewers: weekly price decorrelation between legs is **not** required and usually will not exist. Inversion and complementarity are claimed at the *scenario* level only. Any pair or trio sold to you on daily decorrelation is being oversold.




Fig. 3 — Founding books vs a naive pair (anchor-interpolated, Jul 2025–Jul 2026). The naive gold/silver pair beat the Metals trio +41.5% vs +36.1% — and the silver kill was still upheld, on redundancy and substitution risk. The backtest is evidence, not oracle.

IV • From pairs with an edge to the three-seat book

Pairs with an arbitrage-like relationship were then reviewed across different sectors — energy, metals, agriculture — asking in each case: where is the spread the hypothesis actually predicts, and what regime breaks it? That review produced the project's core structure, the **TRIO**: every book fills exactly three seats, and an asset may qualify for one seat only.

The **Engine** is thesis-central growth, qualified by at least two dated catalysts inside the horizon. The **Tollbooth** collects regardless of which competitor wins — regulated returns, royalties and streams at locked costs, backbone infrastructure — ideally with a counter-cyclical mechanism that strengthens when the sector is weak. This is the Thales test in one line: *presses, not pickaxes*. The **Regime-Inverse** is the leg whose best macro scenario is a headwind for the other two, subject to non-cannibalization: all three must still be able to win together in a sector-wide bull.

Two founding books calibrate the method. **Electrons** — CATL as engine, National Grid as tollbooth, Fluence as regime-inverse on the US–China policy axis, weighted 45/30/25 — backtested +108.4% with an anchor max drawdown of –26.2%; Eos was killed at K3/K4 (–75% drawdown plus a rights issue) and Tesla at K2 (its Shanghai Megafactory broke the policy-purity seat). **Metals** — gold as ballast, platinum as regime-inverse (peace-and-boom is gold's headwind), Wheaton as tollbooth via streams contracted at \$473/oz gold and \$5.75/oz silver, weighted 40/25/35 — backtested +36.1% with –28.9% drawdown, and retained silver exposure elegantly through Wheaton's streams after silver itself was killed. Cross-sector league tables are forbidden: the same calendar means different regimes, and +108% in a supercycle year is not a better construction than +36% in a crash-recovery year.




Fig. 4 — The three seats. Weights default to Engine 40–45%, Tollbooth 30–35%, Regime-Inverse 25–30%; no leg ever exceeds 60%.

V • The Cicada Hunt — scheduled black swans

Finally, the **Cicada Hunt** was established to identify verifiable events the market may not yet have priced correctly — or priced at all. A cicada is a rare, high-impact event whose **date is already fixed** by a biological, astronomical, legal, software-epoch, demographic, or infrastructure clock. Five gates must all pass: the event is **scheduled** (probabilistic tails like hurricanes fail — those are swans, not cicadas); its recurrence exceeds roughly seven years, so **institutional amnesia** is the edge; the **transmission** is mapped with sesame-bun specificity (who bleeds, who feeds, mechanism, lag); the **obscurity test** passes — saturation in trade press with silence in financial coverage is the classic signature, while a sell-side preview note within twelve months means it is priced, and killed; and a **precedent audit** of the last emergence justifies, in writing, the probability used to underwrite it.

Survivors become small asymmetric paper tickets under an expected-value floor of +0.30 per unit staked, with $EV = p \cdot (M-1) - (1-p) \cdot L$, tiered instruments (options underwritten at 10x minimum payoff, pure-play equities at 3x), a hard 5% sleeve cap, and a breadth doctrine — many small tickets beat one big one. Every ticket self-liquidates on its event date. The live example: **CH-001**, the 2 August 2027 total eclipse over Egypt (6m23s of totality, the longest on accessible land until 2114), Tier B at $p=0.35$, $M=3 \rightarrow EV +0.41$, entered on evidence of fully prepaid eclipse-week bookings a year in advance while financial coverage remained silent.




Fig. 5 — The five gates, and the EV floor by instrument tier.

VI • Summary and standing frame

The playbook, in one paragraph: write the hypothesis first; scour for tickers second; let pre-registered kills execute themselves; prefer three complementary seats to one loud bet; buy presses, not pickaxes; hunt events with fixed dates and forgotten precedents; grade the mechanism, not the luck — a green P&L; with a dead mechanism is a rejection that got lucky. All of it is **paper trading only**: a hypothetical learning ledger, not financial advice and not a solicitation. Nothing here has been executed with real money, backtests rest on coarse anchors from a single regime of history with drawdowns understated, and reviewers are invited to attack the kill criteria hardest of all — that is where the method lives or dies.